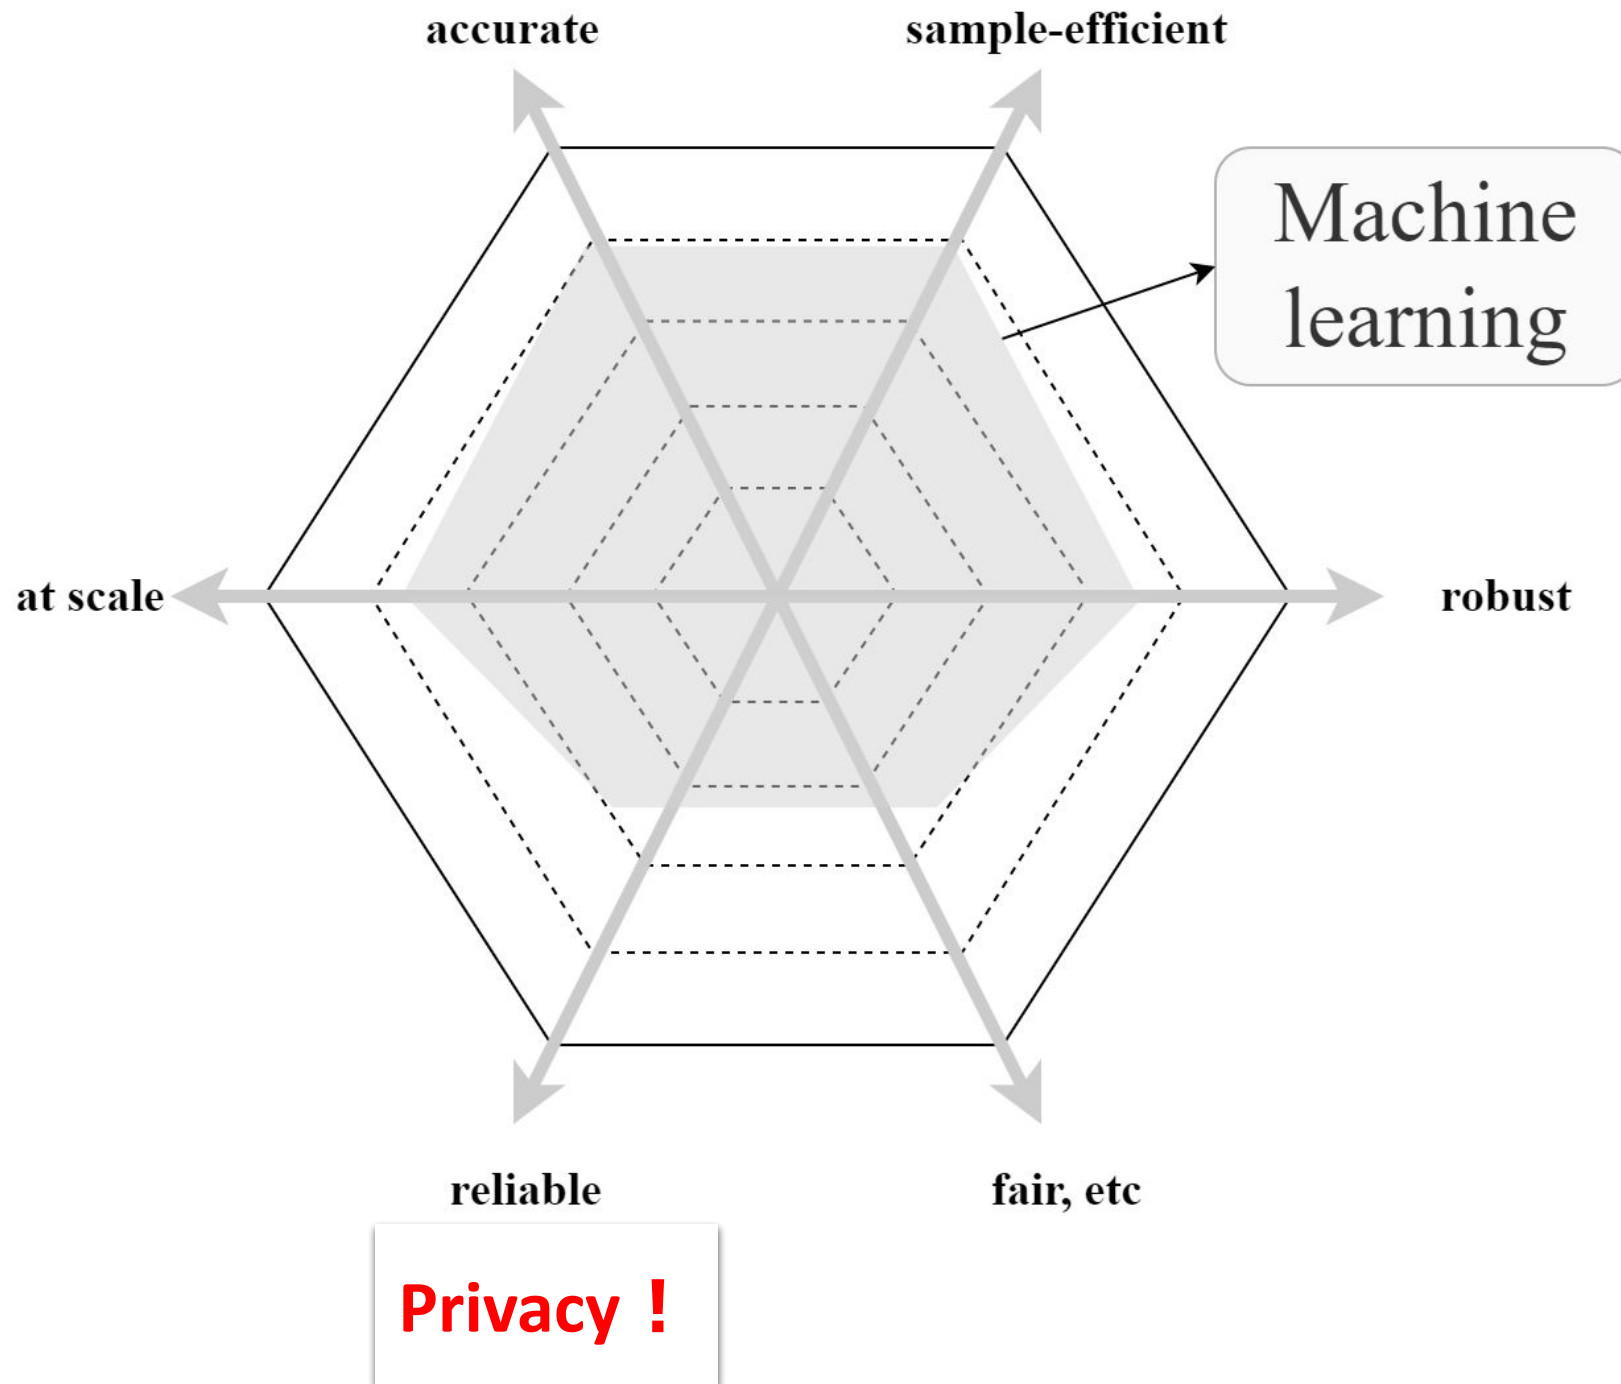


Reconstructing Training Data from Model Gradient, Provably

Qi Lei, Courant Math and CDS

With Zihan Wang, Jason Lee

<https://arxiv.org/abs/2212.03714>

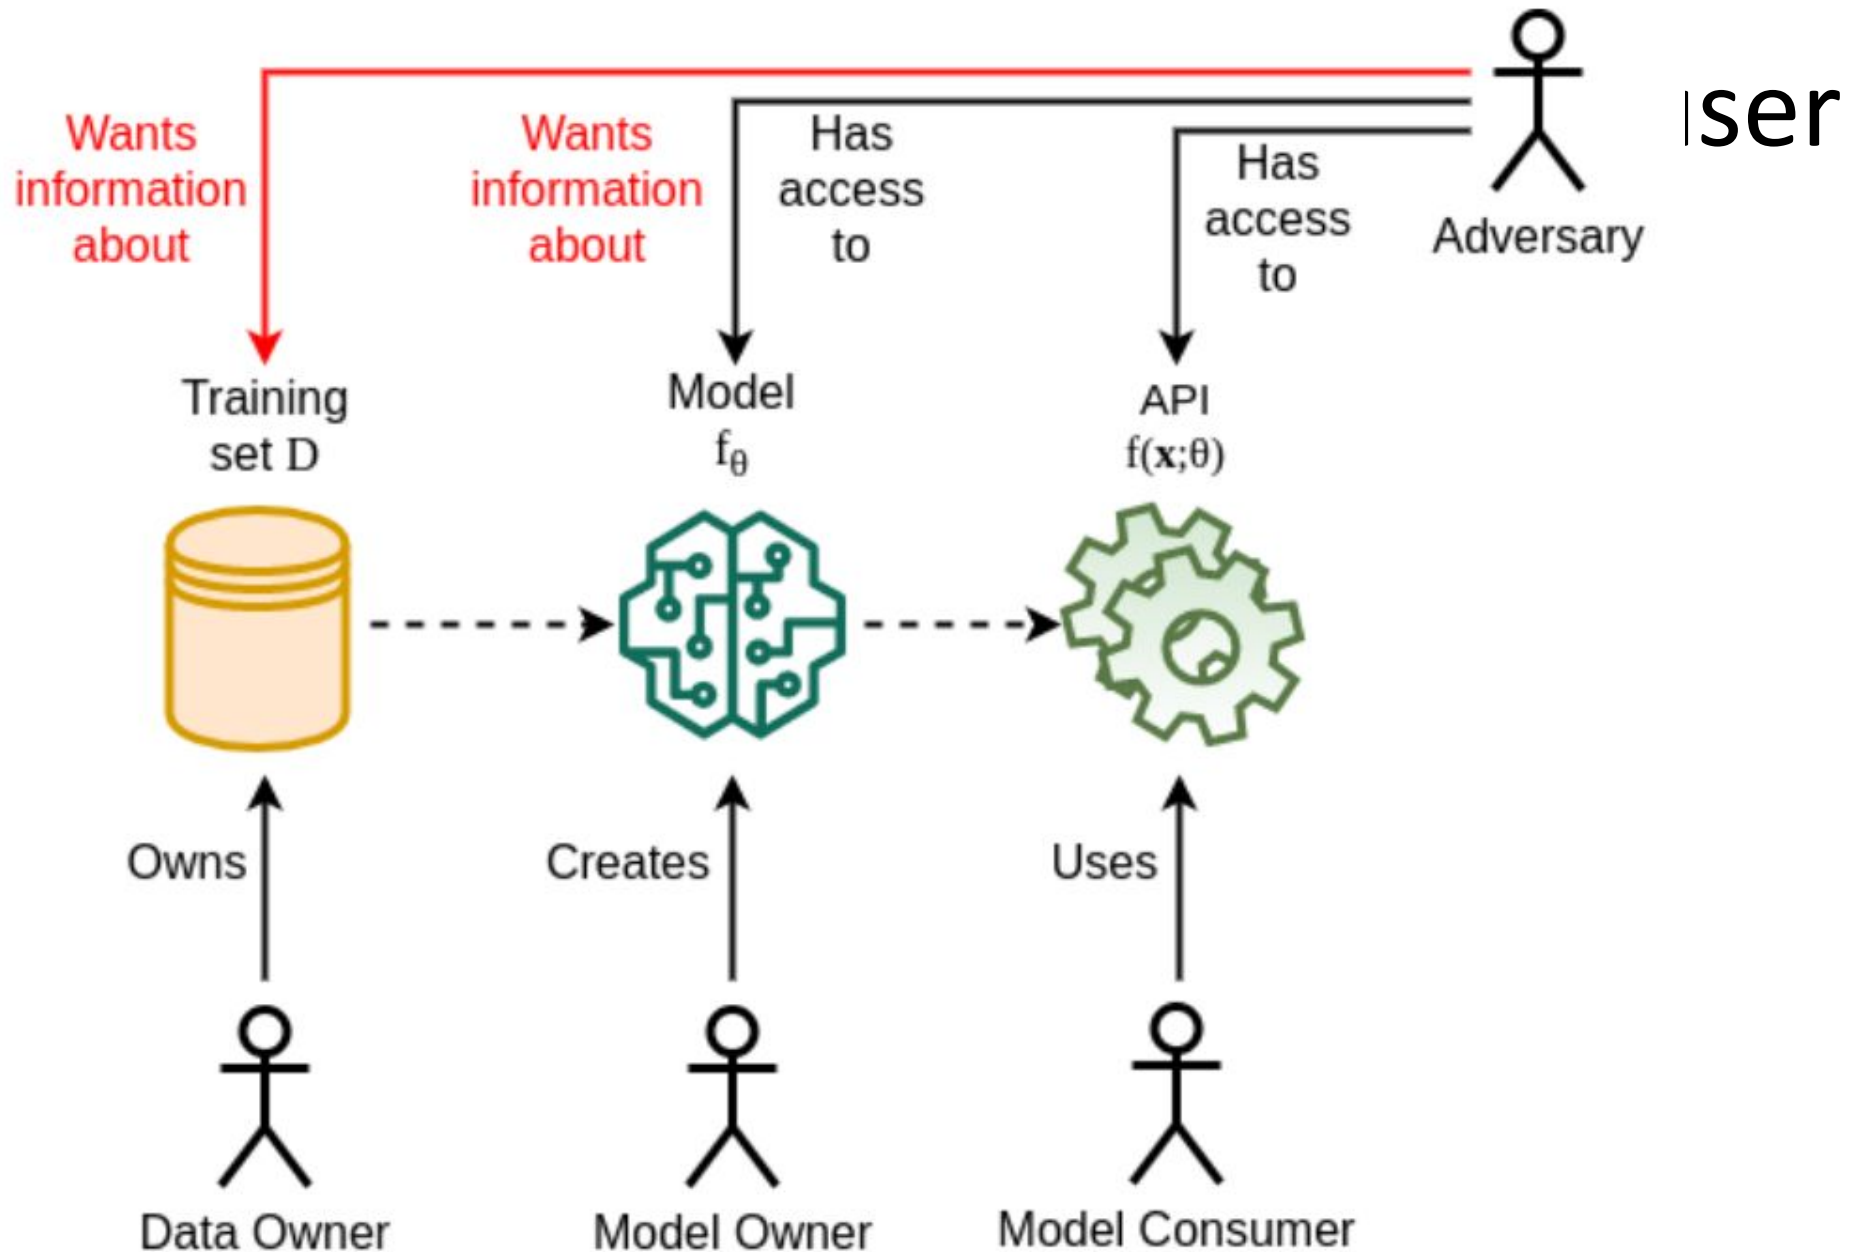


Privacy attacks: infer information about user data/protected model

- Data curation stage
 - Data linkage attack (reidentification)
- Model training phase
 - Data reconstruction attack
- Inference/prediction time
 - Membership inference attack (tracing attack)
 - Data reconstruction attack

Private data

- Data collection
- Model training
- Inference



Resemblance to inverse problems

Different types of privacy attacks

- Membership Inference Attacks (differential privacy) [Shokri et al 2017]
- Reconstruction Attacks (our target) [Gong&Niu,2016]
- Property Inference Attacks [Ateniese et al 2015]
- Model Extraction Attacks (knowledge distillation) [Papernot et al 2018]

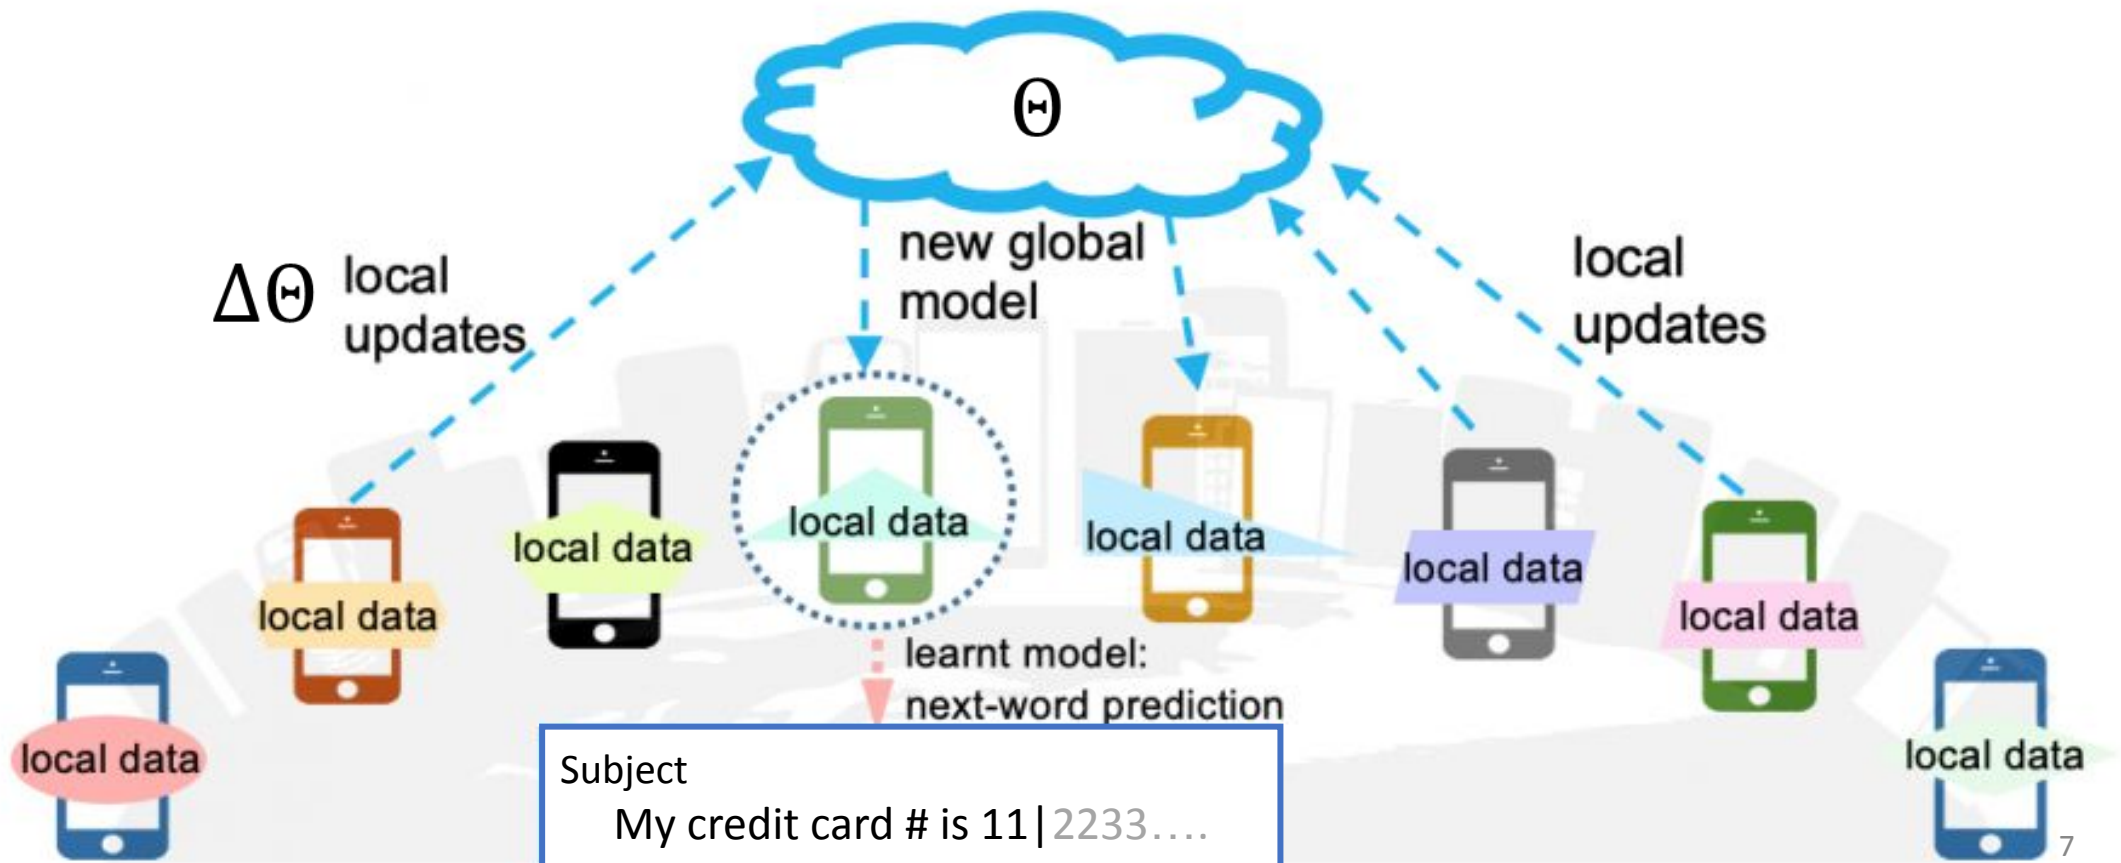
Different types of privacy attacks

- Membership Inference Attacks (differential privacy) [Shokri et al 2017]
- **Reconstruction Attacks (our target)** [Gong&Niu,2016]
- Property Inference Attacks [Ateniese et al 2015]
- Model Extraction Attacks (knowledge distillation) [Papernot et al 2018]

Federated learning

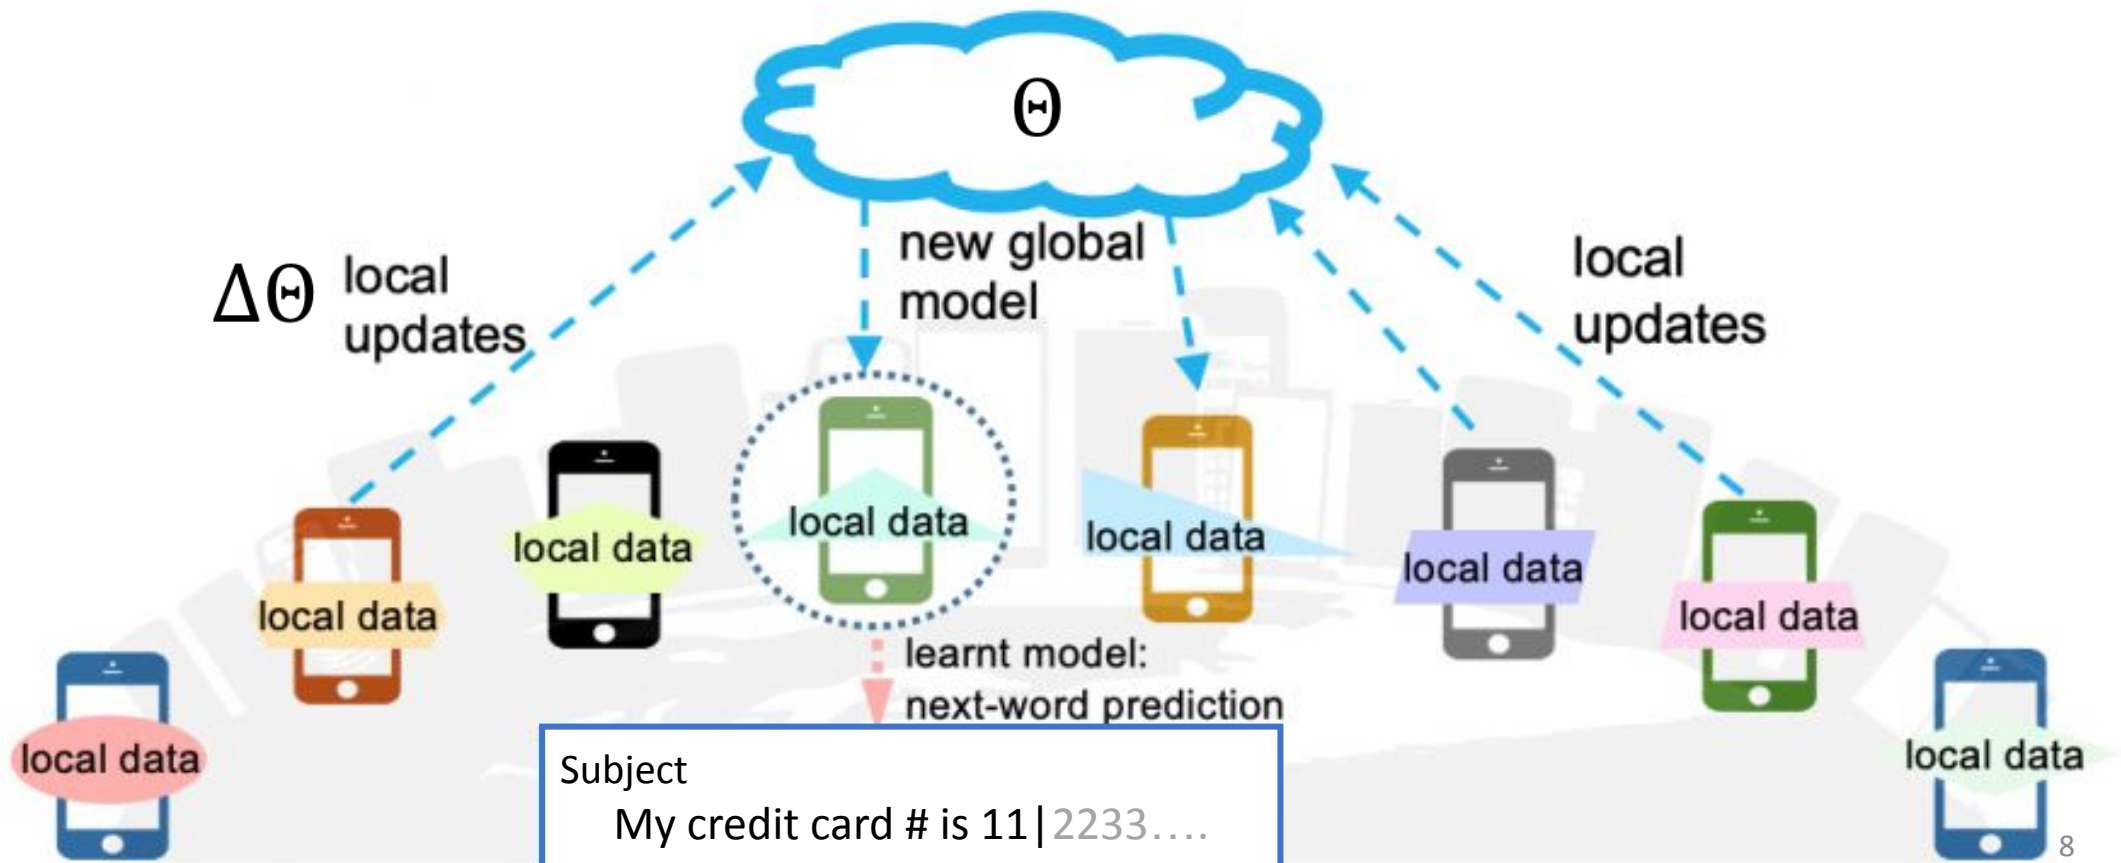
[Konečný et al. 2016, McMahan et al. 2017]

- Privacy leakage in distributed learning - Data and model not co-located



Privacy leakage in distributed learning

- Does local update reveal the training data?



Federated learning



Gautam Kamath
@thegautamkamath



Federated learning, wherein training data never leaves the user's device (only gradients or model parameters), is an effective way to protect the privacy of individual training points.

True

19.4%

False

46.1%

I don't know/show me

34.5%

956 votes · Final results

11:32 PM · Dec 11, 2022

More formal statement

- Local batch of data: $S = \{(x_1, y_1), (x_2, y_2), \dots, (x_B, y_B)\}$
- Prediction function: $x \rightarrow f(x; \Theta)$
- Local update: $G := \frac{1}{B} \nabla_{\Theta} \sum_{i=1}^B \ell(f(x_i, \Theta), y_i)$

Fundamental questions

- Is the model gradient G sufficient to identify the training samples?
- If so, is there efficient algorithm to recover the samples?

Prior work

- Attacking methods

- Learn to generate the training samples from a local user
- Match the gradient: $\min_{S=\{(x_i, y_i)\}} \left\| G - \sum_{i=1}^B \nabla \ell(f(x_i; \Theta), y_i) \right\|^2$
- Model inversion at inference time

- Defending methods

- Quantizing the gradient
- Add noise

Some folklore from empirical findings

What parameters to query at?

- Moderately trained model
 - Random network hasn't memorized the data,
 - Well-trained model makes gradient vanish

Wrong
impressions!

Is gradient alone enough to identify the images?

- Prior work believed not.
 - Introduce prior information of the training data (model by generative models)

Setting

- Warm-up: two-layer neural network

$$f(x; \{W, a\}) = \sum_{j=1}^m a_j \sigma(w_j^\top x) = a^\top \sigma(W^\top x)$$

- Choose proper w_j, a_j to query the gradient at

$$\nabla_{a_j} L = \sum_{i=1}^B l'_i \sigma(w_j^\top x_i)$$

Caveat on linear and quadratic activations:

- Linear setting:
- $\nabla_a L = W \left(\sum_{i=1}^B l'_i x_i \right); \nabla_W L = a \left(\sum_{i=1}^B l'_i x_i \right)^\top$
- Can only identify a linear combination of X

Caveat on linear and quadratic activations:

- Linear setting:
 - $\nabla_a L = W \left(\sum_{i=1}^B l'_i x_i \right); \nabla_W L = a \left(\sum_{i=1}^B l'_i x_i \right)^\top$
 - Can only identify a linear combination of X
- Quadratic setting:
 - $\nabla_{a_j} L = w_j^\top \bar{\Sigma} w_j; \nabla_{w_j} L = 2\bar{\Sigma} w_j$, here $\bar{\Sigma} = \sum_{i=1}^B l'_i x_i x_i^\top$
 - Can only identify the span of X

Our algorithm: using Stein's lemma

- Stein's lemma: $E_{w \sim N(\mathbb{0}, I)} [g(a^\top w) H_p(w)] = E[g^{(p)} a^{\otimes p}]$.
- Hermite function: $H_2(w) = ww^\top - I, H_3(w) = w^{\otimes 3} - w \widetilde{\otimes} I$.
- $\widehat{T}_p := \frac{1}{m} \sum_{j=1}^m g(w_j) H_p(w_j) \approx E_{w \sim N(\mathbb{0}, I)} [g(w) H_p(w)]$
 $\equiv \sum_{i=1}^B E \left[\sigma^{(p)}(w^\top x_i) x_i^{\otimes p} \right] =: T_p$
- $g(w_j) := \nabla_{a_j} L = \sum_{i=1}^B l'_i \sigma(w_j^\top x_i)$ is our observation from the model gradient

Tensor decomposition

- Now we can estimate $T_p := \sum_{i=1}^B E \left[\sigma^{(p)}(w^\top x_i) x_i^{\otimes p} \right]$
- Uniquely identify $\{x_1, x_2, \dots, x_B\}$ through tensor decomposition when data is linearly independent for $p \geq 3$. [Kuleshov et al. 2015]
- Our strategy: choose $a_j = \frac{1}{m}$, $w_j \sim N(0, I)$, estimate T by $\widehat{T}_3 := \frac{1}{m} \sum_{j=1}^m g(w_j) H_3(w_j)$, $g(w_j) := \nabla_{a_j} L$

Improving the dimension dependence

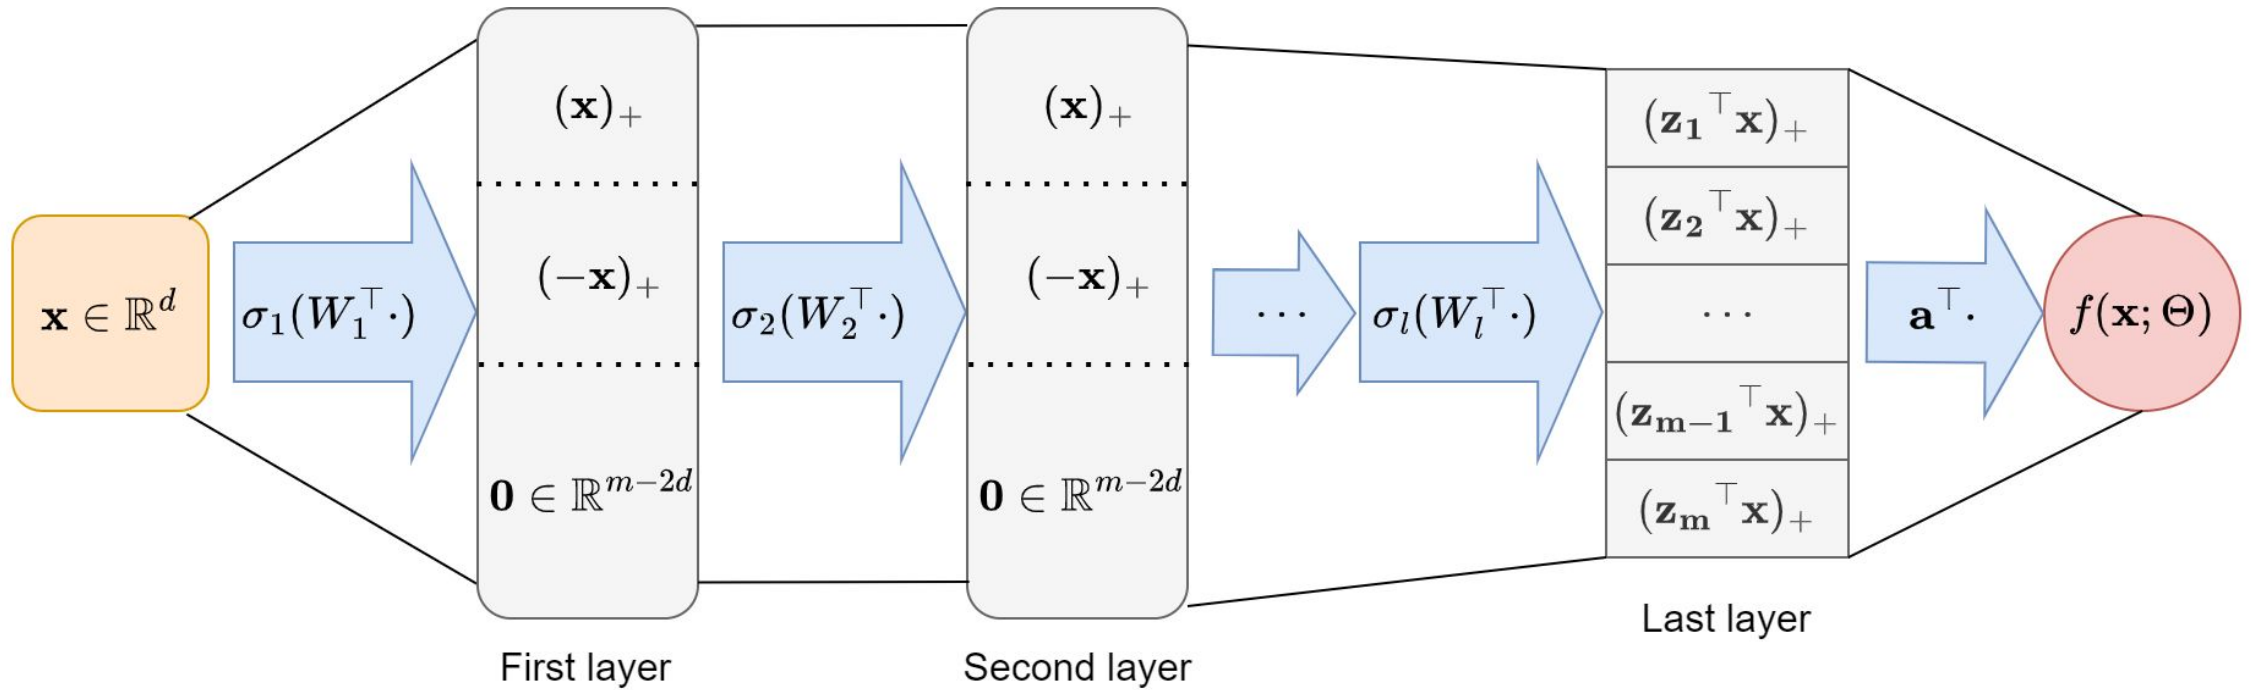
- First estimate the span of $\{x_1, x_2, \dots, x_B\}$ by decomposing $T_2 = VV^\top$. $V \in R^{d \times B}$.
- Find $T_3(V, V, V) \in R^{B \times B \times B}$ and conduct tensor decomposition $\{u_1, u_2, \dots, u_B\}$.
- Project back to the original space $\hat{x}_j = Vu_j$.
- Can also use the estimated x as initialization and do gradient matching.

-relevant strategy appeared in [Zhong et al 2017] for optimizing over 2-layer neural network (dual problem)

Theoretical analysis

- Applies when $E[\sigma^{(3)}(w)]$ or $E[\sigma^{(4)}(w)] \neq 0$. Applies to sigmoid, tanh, ReLU, leaky ReLU.
- Reconstruction error $\leq \tilde{O}(\sqrt{d/m})$.

Extension to deeper neural networks



- Previous findings: if last two layers are fully connected, can recover the features from the $(l - 2)$ -th layer
- Other structured data modalities: recover the embeddings first

Discussions

- Identifiability
 - $E[\sigma^{(3)}(w)]$ or $E[\sigma^{(4)}(w)] \neq 0$ (work for most activation functions)
 - # of hidden nodes m scales at least linearly with d
 - Deeper neural network does not help or hurt
- Distinction to linear case or convex optimization:
 - Linear and neural networks are fundamentally different on whether the gradient leaks data
- Inspiration on defending algorithms:
 - Adding noise does not help

Thank you