

Lei, Qi

Website: <https://cecilialeiqi.github.io/>

Google Scholar: <https://scholar.google.com/citations?user=kG0gaowAAAAJ&hl=en>

Email: qilei@princeton.edu

Professional Experience

- New York University, NY, United States* September 2022 -
• Assistant Professor in Mathematics and Data Science, Courant Institute of Mathematical Sciences and the Center for Data Science
- Princeton University, NJ, United States* July 2020 - August 2022
• Associate Research Scholar and Postdoc Research Associate (CIFellow), Electrical and Computer Engineering Department
• Mentor: Jason D. Lee
- Institute for Advanced Study, Princeton, NJ, United States* September 2019 - July 2020
• Visiting Graduate Student for the “Special Year on Optimization, Statistics, and Theoretical Machine Learning”
- Simons Institute, Berkeley, CA, United States* May 2019 - August 2019
• Research Fellow for the Foundations of Deep Learning Program
- Amazon/A9 Product Search* May 2017 - August 2017
• Inline search suggestions
- Amazon Web Services (AWS Deep Learning Team)* January 2017 - April 2017
• Deep Learning Tutorial

Education

- University of Texas at Austin, TX, United States* August 2014 - May 2020
• Ph.D., Oden Institute for Computational Sciences and Engineering
• Advisors: Alexandros G. Dimakis and Inderjit S. Dhillon

Selected Publications

(* indicates α - β order)

1. Jianwei Li, **Qi Lei**, Wei Cheng, Dongkuan Xu. “Towards Robust Pruning: An Adaptive Knowledge-Retention Pruning Strategy for Language Models”, *To appear at EMNLP conference, 2023*
2. Zihan Wang, Jason Lee, **Qi Lei**. “Reconstructing Training Data from Model Gradient, Provably”, *AISTATS 2023: 6595-6612*
3. Jason D. Lee*, **Qi Lei***, Nikunj Saunshi*, Jiacheng Zhuo*, “Predicting What You Already Know Helps: Provable Self-Supervised Learning”, *NeurIPS 2021: 309-323*
4. Simon S. Du*, Wei Hu*, Sham M. Kakade*, Jason D. Lee*, **Qi Lei***. “Few-Shot Learning via Learning the Representation, Provably”, *ICLR, 2021*
5. **Qi Lei**, Ajil Jalal, Inderjit S. Dhillon, Alexandros G. Dimakis. “Inverting Deep Generative models, One layer at a time”, *NeurIPS 2019: 13910-13919*
6. **Qi Lei**, Lingfei Wu, Pin-Yu Chen, Alexandros G. Dimakis, Inderjit S. Dhillon, Michael Witbrock. “Discrete Adversarial Attacks and Submodular Optimization with Applications to Text Classification”, *MLSys 2019 (covered by Nature News)*